

```
import numpy as np
import keras
from keras import layers

# Model / data parameters
num_classes = 10
input_shape = (28, 28, 1)

# Load the data and split it between train and test sets
(x_train, y_train), (x_test, y_test) = keras.datasets.mnist.load_data()

# Scale images to the [0, 1] range
x_train = x_train.astype("float32") / 255
x_test = x_test.astype("float32") / 255
# Make sure images have shape (28, 28, 1)
x_train = np.expand_dims(x_train, -1)
x_test = np.expand_dims(x_test, -1)
print("x_train shape:", x_train.shape)
print(x_train.shape[0], "train samples")
print(x_test.shape[0], "test samples")

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

# Build the model
model = keras.Sequential(
    [
        keras.Input(shape=input_shape),
        layers.Conv2D(32, kernel_size=(3, 3), activation="relu"),
        layers.MaxPooling2D(pool_size=(2, 2)),
        layers.Conv2D(64, kernel_size=(3, 3), activation="relu"),
        layers.MaxPooling2D(pool_size=(2, 2)),
        layers.Flatten(),
        layers.Dropout(0.5),
        layers.Dense(num_classes, activation="softmax"),
    ]
)

model.summary()

# Train the model
batch_size = 128
epochs = 15

model.compile(loss="categorical_crossentropy", optimizer="adam", metrics=["accuracy"])

model.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, validation_split=0.1)

# Evaluate the trained model
score = model.evaluate(x_test, y_test, verbose=0)
print("Test loss:", score[0])
print("Test accuracy:", score[1])
```

 Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>  
 11490434/11490434 1s 0us/step  
 x\_train shape: (60000, 28, 28, 1)  
 60000 train samples  
 10000 test samples  
 Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d ( <a href="#">Conv2D</a> )	(None, 26, 26, 32)	320
max_pooling2d ( <a href="#">MaxPooling2D</a> )	(None, 13, 13, 32)	0
conv2d_1 ( <a href="#">Conv2D</a> )	(None, 11, 11, 64)	18,496
max_pooling2d_1 ( <a href="#">MaxPooling2D</a> )	(None, 5, 5, 64)	0
flatten ( <a href="#">Flatten</a> )	(None, 1600)	0
dropout ( <a href="#">Dropout</a> )	(None, 1600)	0
dense ( <a href="#">Dense</a> )	(None, 10)	16,010

Total params: 34,826 (136.04 KB)

Trainable params: 34,826 (136.04 KB)

Non-trainable params: 0 (0.00 B)

Epoch 1/15

422/422 59s 134ms/step - accuracy: 0.7670 - loss: 0.7613 - val\_accuracy: 0.9772 - val\_loss: 0.0847

Epoch 2/15

422/422 73s 113ms/step - accuracy: 0.9633 - loss: 0.1228 - val\_accuracy: 0.9848 - val\_loss: 0.0586

Epoch 3/15

422/422 46s 110ms/step - accuracy: 0.9731 - loss: 0.0904 - val\_accuracy: 0.9880 - val\_loss: 0.0454

Epoch 4/15

422/422 83s 112ms/step - accuracy: 0.9776 - loss: 0.0706 - val\_accuracy: 0.9867 - val\_loss: 0.0446

Epoch 5/15

422/422 81s 111ms/step - accuracy: 0.9798 - loss: 0.0664 - val\_accuracy: 0.9895 - val\_loss: 0.0399

Epoch 6/15

422/422 81s 109ms/step - accuracy: 0.9827 - loss: 0.0551 - val\_accuracy: 0.9892 - val\_loss: 0.0378

Epoch 7/15

422/422 49s 117ms/step - accuracy: 0.9839 - loss: 0.0531 - val\_accuracy: 0.9893 - val\_loss: 0.0380

Epoch 8/15

422/422 46s 110ms/step - accuracy: 0.9853 - loss: 0.0480 - val\_accuracy: 0.9900 - val\_loss: 0.0349

Epoch 9/15

422/422 82s 111ms/step - accuracy: 0.9850 - loss: 0.0475 - val\_accuracy: 0.9902 - val\_loss: 0.0326

Epoch 10/15

422/422 84s 117ms/step - accuracy: 0.9865 - loss: 0.0452 - val\_accuracy: 0.9912 - val\_loss: 0.0333

Epoch 11/15

422/422 81s 115ms/step - accuracy: 0.9879 - loss: 0.0391 - val\_accuracy: 0.9910 - val\_loss: 0.0330

Epoch 12/15

422/422 82s 114ms/step - accuracy: 0.9860 - loss: 0.0429 - val\_accuracy: 0.9903 - val\_loss: 0.0325

Epoch 13/15

422/422 47s 112ms/step - accuracy: 0.9884 - loss: 0.0370 - val\_accuracy: 0.9922 - val\_loss: 0.0302

Epoch 14/15

422/422 81s 111ms/step - accuracy: 0.9893 - loss: 0.0343 - val\_accuracy: 0.9922 - val\_loss: 0.0286

Epoch 15/15

422/422 84s 116ms/step - accuracy: 0.9877 - loss: 0.0373 - val\_accuracy: 0.9917 - val\_loss: 0.0304

Test loss: 0.025226449593901634